

# An intro to Boost.PolyCollection



using `std::cpp` 2017

Joaquín M López Muñoz <joaquin.lopezmunoz@gmail.com>

Madrid, November 2017



Before the show begins: some news from the Boost front



# Before the show begins: some news from the Boost front

- 145 libraries (1.65.1) and counting, 5 new last year
  - Fiber, QVM, Process, PolyCollection, Stacktrace
- Trends
  - C++11 or later increasingly required by new libs
  - Postmodern (meta)functional: Hana, Fit, Mp11
  - OS and communications: Beast, Fiber, Process, Stacktrace
- Building Boost
  - Boost.Build challenged by CMake
  - Bincrafters now provides Boost packages through **conan**
- You have no excuses for not using Boost!



Let 's dive into Boost.PolyCollection



# Type hierarchy for a role playing game

```
struct sprite
{
    sprite(int id):id(id){}
    virtual ~sprite()=default;
    virtual void render(std::ostream& os) const=0;

    int id;
};
```

```
struct warrior:sprite
{
    using sprite::sprite;
    warrior(std::string rank,int id):
        sprite{id},rank{std::move(rank)}{}
    void render(std::ostream& os) const override
    {os<<rank<<" "<<id;}
    std::string rank="warrior";
};
```



```
struct juggernaut:warrior
{
    juggernaut(int id):
        warrior{"juggernaut",id}{}
};
```



```
struct goblin:sprite
{
    using sprite::sprite;
    void render(std::ostream& os)
    const override
    {os<<"goblin "<<id;}
};
```



# Populate and render



vector.cpp

```
std::vector<std::unique_ptr<sprite>> c;

std::mt19937 gen{92748}; // some arbitrary random seed
std::discrete_distribution<> rnd{{1,1,1}};
for(int i=0;i<8;++i){ // assign each type with 1/3 probability
    switch(rnd(gen)){
        // watch out: std::make_unique requires C++14
        case 0: c.push_back(std::make_unique<warrior>(i));break;
        case 1: c.push_back(std::make_unique<juggernaut>(i));break;
        case 2: c.push_back(std::make_unique<goblin>(i));break;
    }
}

const char* comma="";
for(const auto& p:c){
    std::cout<<comma;
    p->render(std::cout);
    comma=",";
}
std::cout<<"\n";
```

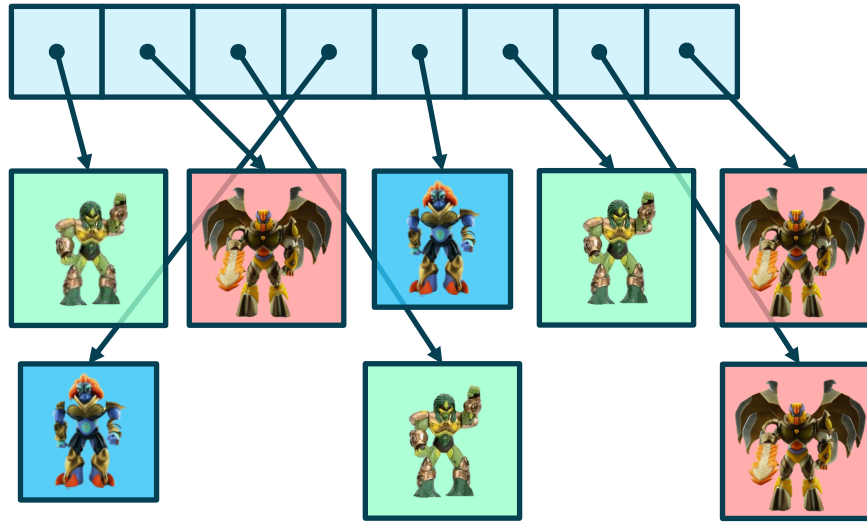


# Piece of cake



GORMITI is a trademark of Giochi Preziosi S.p.A.

# There are hidden inefficiencies, though

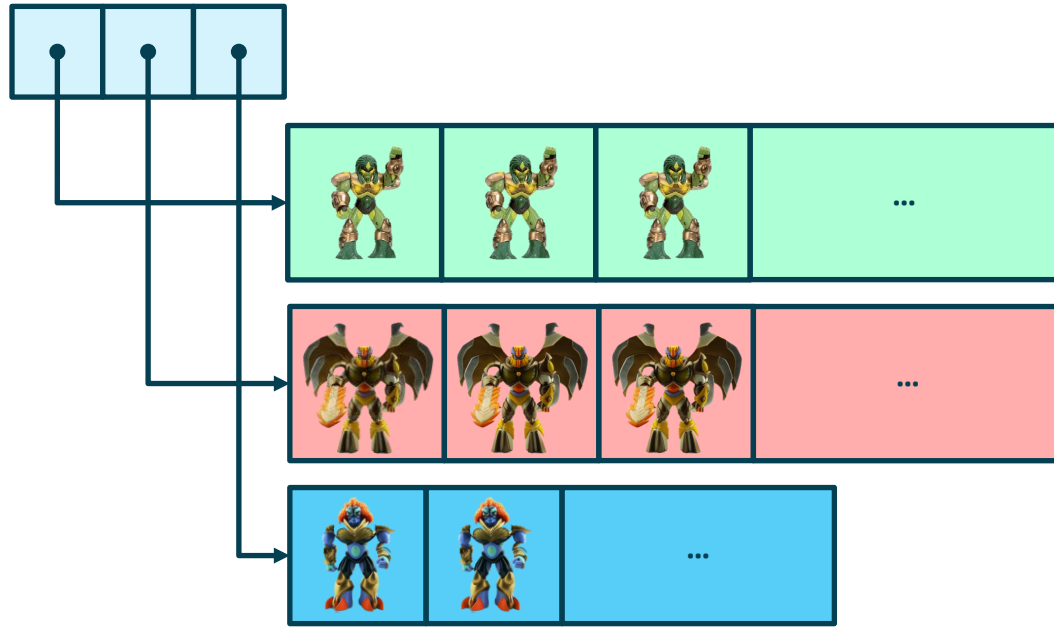


- Objects are scattered through memory → CPU cache misses
- Types are mixed up along traversal → CPU branch prediction fails
  - We could sort the vector on type\*, object scattering is harder to fix though
- All boils down to OOP requiring pointer indirections

\*See my “Mind the cache” talk at [github.com/joaquintides/usingstdcpp2015](https://github.com/joaquintides/usingstdcpp2015)



# Packing things together



- `boost::base_collection` stores elements in per-type dedicated **segments**
- `[begin(), end())` traverses each segment in succession
  - So, free ordering of elements is lost (except within segment)
- In return, performance increases spectacularly



```
boost::base_collection<sprite> c;

std::mt19937 gen{92748}; // some arbitrary random seed
std::discrete_distribution<> rnd{{1,1,1}};
for(int i=0;i<8;++i){ // assign each type with 1/3 probability
    switch(rnd(gen)){
        case 0: c.insert(warrior{i});break;
        case 1: c.insert(juggernaut{i});break;
        case 2: c.insert(goblin{i});break;
    }
}

const char* comma="";
for(const sprite& s:c){
    std::cout<<comma;
    s.render(std::cout);
    comma=", ";
}
std::cout<<"\n";
```

- Spot the differences with the `std::vector` code



# Interface at a glance



`base_collection_interface.cpp`

**Segments can be targeted individually**

**Insertion interface mimics `std::multiset` rather than `std::vector`**

**Emplacement requires that type of element be specified**

**New segments are created automatically; otherwise, use `register_types`**



## Segments can be targeted individually

```
std::sort( // sort warriors in descending id order
    c.begin<warrior>(), c.end<warrior>(),
    [](const warrior& x, const warrior& y){return x.id > y.id;});

for(const warrior& w: c.segment<warrior>()) // print warrior ids
    std::cout << w.id << "\n";

c.erase(c.begin(typeid(warrior))); // erase first warrior

c.clear<warrior>(); // erase all warriors
```

- Watch out: `begin(typeid(warrior))` is not the same as `begin<warrior>()`
- Two kinds of local iterators (prefer the latter)

Insertion interface mimics `std::multiset` rather than `std::vector`

Emplacement requires that type of element be specified

New segments are created automatically; otherwise, use `register_types`



# Interface at a glance



base\_collection\_interface.cpp

Segments can be targeted individually

**Insertion interface mimics `std::multiset` rather than `std::vector`**

```
c.insert(juggernaut{8}); // at the end of juggernaut segment  
  
c.insert(c.begin(), juggernaut{9}); // hint only useful if  
                                   // juggernauts are the first segment  
c.insert(c.begin<juggernaut>(), juggernaut{10}); // 1st of juggernauts
```

Emplacement requires that type of element be specified

New segments are created automagically; otherwise, use `register_types`

# Interface at a glance



base\_collection\_interface.cpp

Segments can be targeted individually

Insertion interface mimics `std::multiset` rather than `std::vector`

**Emplacement requires that type of element be specified**

```
c.emplace<goblin>(11);
```

```
c.emplace_pos<goblin>(c.begin<goblin>(),12); // 1st of goblins
```

New segments are created automagically; otherwise, use `register_types`



# Interface at a glance



base\_collection\_interface.cpp

Segments can be targeted individually

Insertion interface mimics `std::multiset` rather than `std::vector`

Emplacement requires that type of element be specified

**New segments are created automatically; otherwise, use `register_types`**

```
// new assortment of sprites
std::array<std::unique_ptr<sprite>,3> a{{
    std::make_unique<elf>(13),
    std::make_unique<ghoul>(14),
    std::make_unique<amazon>(15)
}};

c.register_types<elf,ghoul,amazon>(); // otherwise exception is thrown below

for(const auto&p:a)c.insert(*p);
```

# Interface at a glance



`base_collection_interface.cpp`

**Segments can be targeted individually**

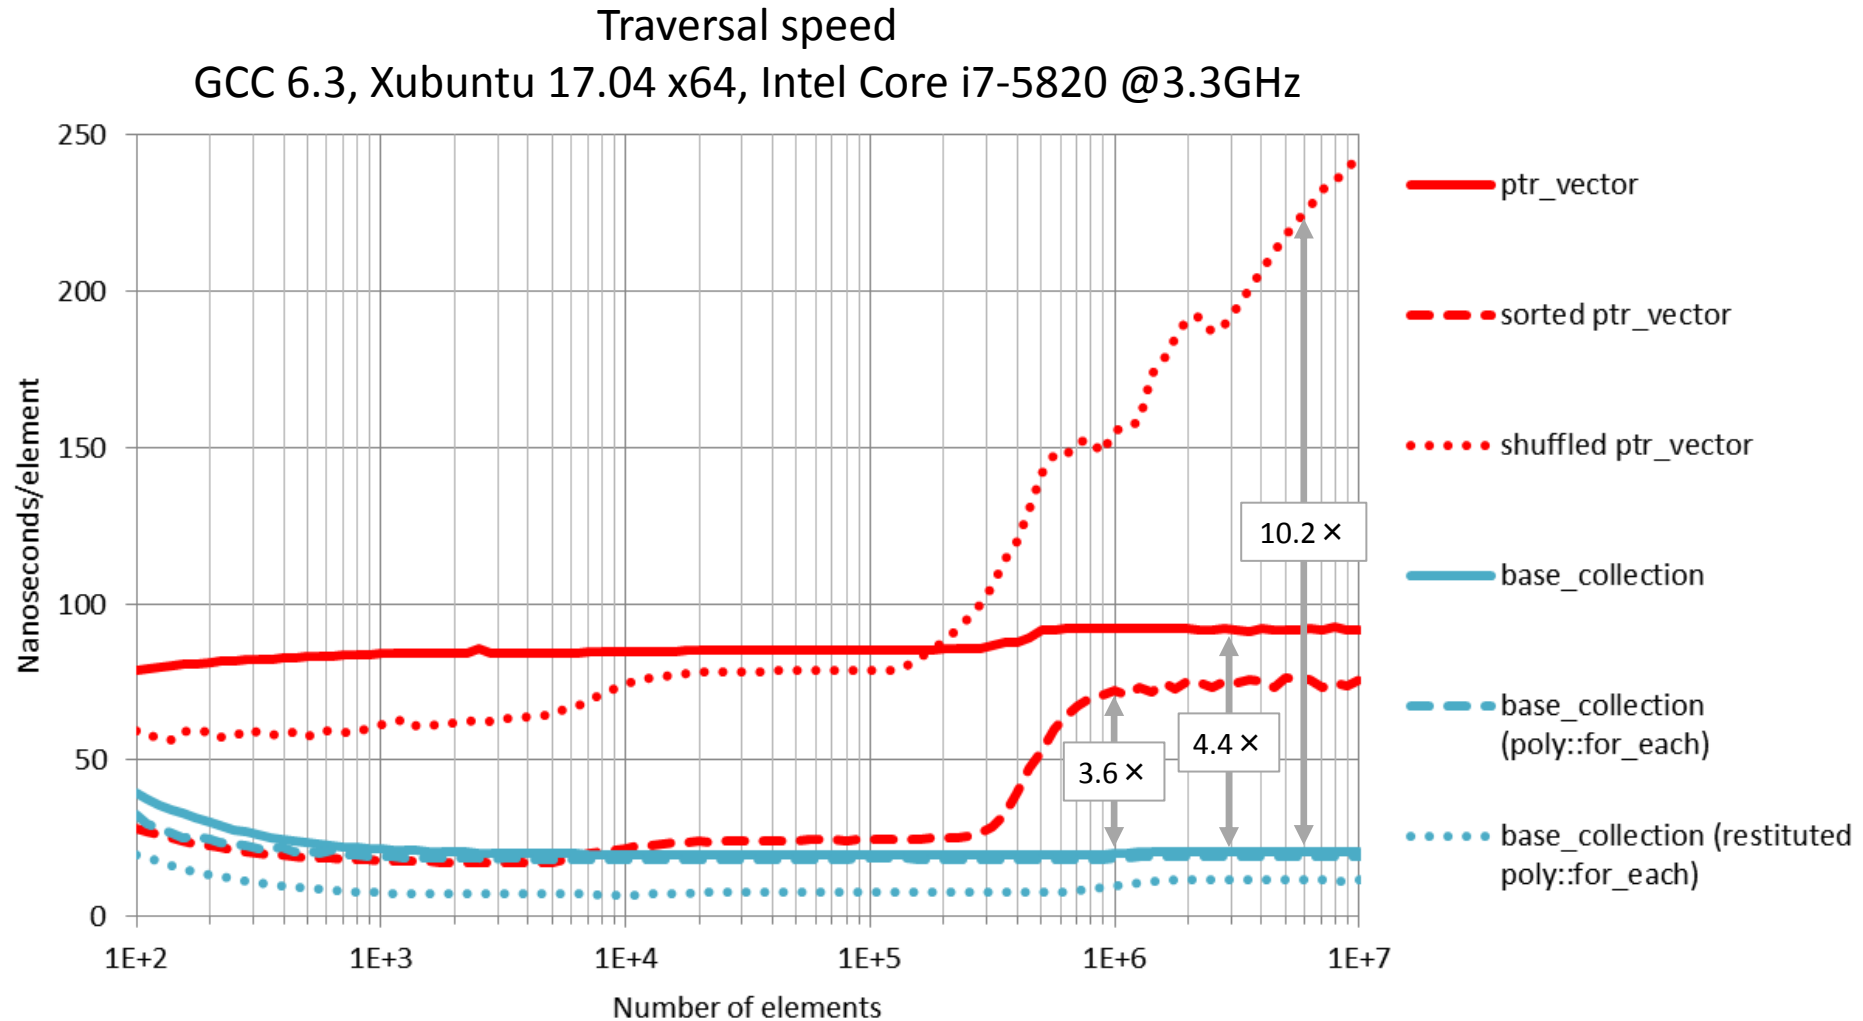
**Insertion interface mimics `std::multiset` rather than `std::vector`**

**Emplacement requires that type of element be specified**

**New segments are created automatically; otherwise, use `register_types`**

- All of this is well, but is it really faster?


# Performance



■ Full results at  
[boost.org/doc/html/poly\\_collection/performance.html](http://boost.org/doc/html/poly_collection/performance.html)



# Performance



■ Full results at  
[boost.org/doc/html/poly\\_collection/performance.html](http://boost.org/doc/html/poly_collection/performance.html)



## ■ Which is fastest?

```
for(const sprite& s:c){  
    std::cout<<comma;  
    s.render(std::cout);  
    comma=",";  
}
```

```
for(const auto& seg_info:c.segment_traversal()){  
    for(const sprite& s:seg_info){  
        std::cout<<comma;  
        s.render(std::cout);  
        comma=",";  
    }  
}
```

## ■ Double loop the basis of Boost.PolyCollection dedicated algorithms

```
boost::poly_collection::for_each(c.begin(),c.end(),[&](const sprite& s){  
    std::cout<<comma;  
    s.render(std::cout);  
    comma=",";  
});
```

## ■ 31 algorithms adapted from <algorithm>



```
boost::poly_collection::for_each<warrior,juggernaut,goblin>(
    c.begin(),c.end(),[&](const auto& s){
        std::cout<<comma;
        s.render(std::cout);
        comma=",";
    }
);
```


- Note the generic lambda

- **Type restitution** is the inverse of type abstraction

```
warrior* pw=new warrior{16};
sprite* ps=pw; // abstraction
warrior* pw2=static_cast<warrior*>(ps); // restitution
```

- Opportunities for **devirtualization** and inlining → more performance

# A new look on dynamic polymorphism



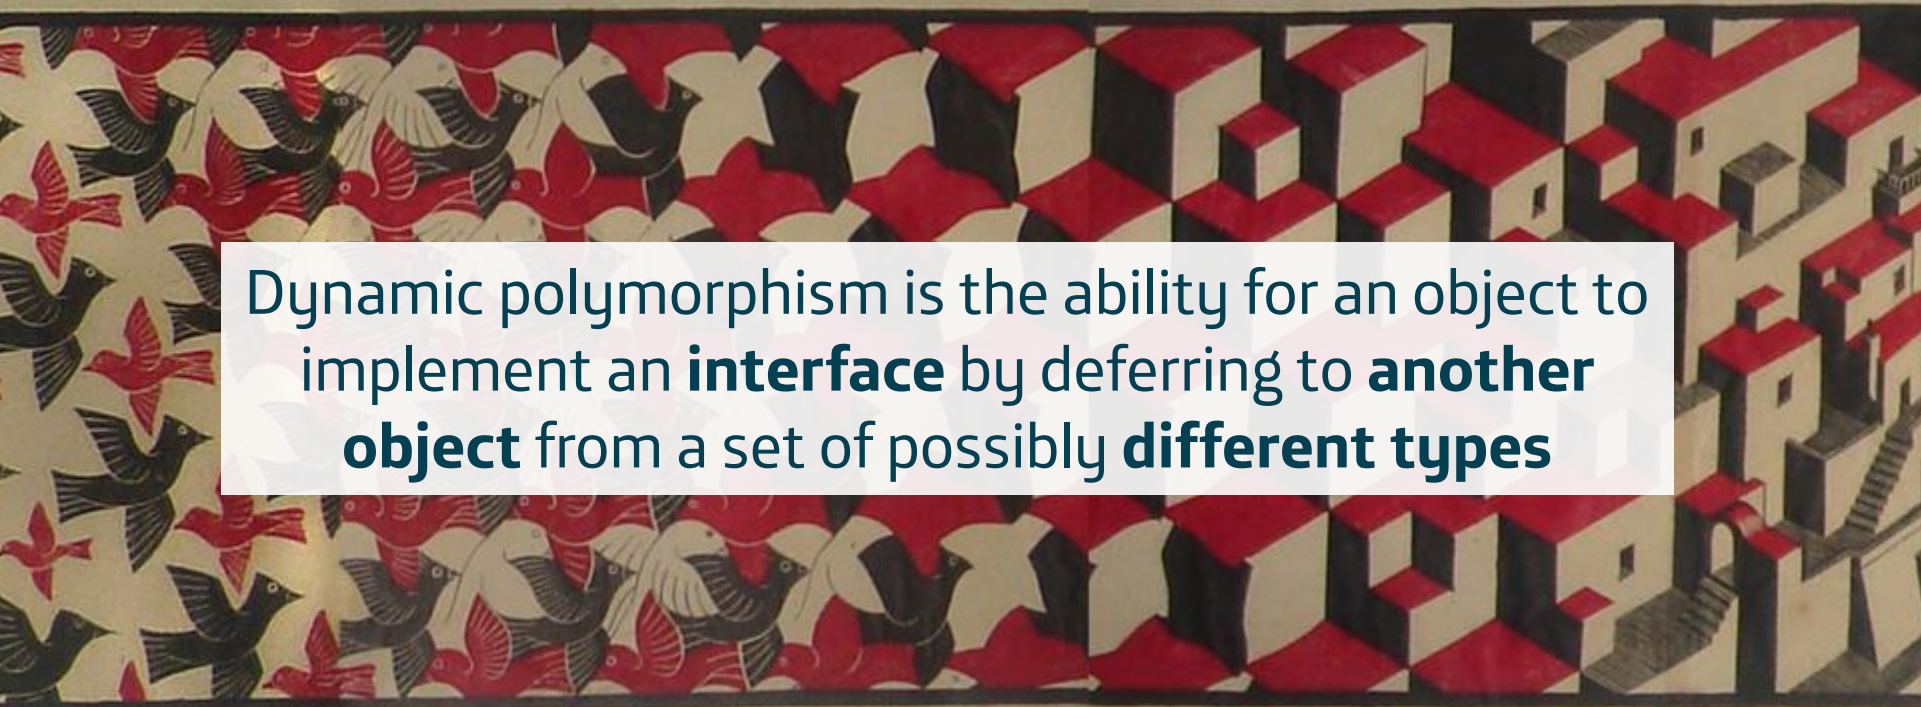


# A new look on dynamic polymorphism

A horizontal strip of a M.C. Escher woodcut, likely from the 'Sky and Water' series. The image shows a transition from a pattern of birds on the left to a pattern of cubes on the right. The birds are rendered in black, white, and red, while the cubes are in shades of red and white. The transition is achieved through a series of intermediate shapes that blend the features of both subjects.

What is dynamic polymorphism, anyway?

# A new look on dynamic polymorphism

A horizontal strip of a M.C. Escher woodcut, likely from the 'Sky and Water' series. The left side features a dense pattern of black birds in flight against a light background. The right side shows a complex, isometric cityscape with red-roofed buildings and a winding path. The central text is overlaid on a white rectangular background.

Dynamic polymorphism is the ability for an object to implement an **interface** by deferring to **another object** from a set of possibly **different types**

# OOP is *just one model* of dynamic polymorphism...

```
sprite(int id);  
virtual ~sprite();
```

```
virtual void render(std::ostream& os) const;
```

interface

```
int id;
```

superobject

```
warrior(std::string rank, int id);
```

```
void render(std::ostream& os) const override;
```

```
std::string rank;
```

subobject

## ...and not a particularly flexible one at that

- Implementation and subtyping are coupled together by inheritance
- Implementation is intentional (again, through inheritance)
- Intersecting interfaces require forethought and virtual inheritance
  - `dynamic_cast` comes into play
  - Ever heard of the diamond problem?\*
  - Believe me, you just don't want to enter this hell

\*Nothing to do with choosing a wedding ring



# Type erasure is the new cool...

- ...and `std::function` its most accessible example

```
int foo(int x){return x;}
struct bar{int operator()(int x)const{return 2*x;}};

using signature=int(int);
std::function<signature> f=&foo,g=bar{},h=[](int x){return 3*x;};

std::cout<<f(1)+g(2)+h(3)<<"\n";
```

- This is a new model of dynamic polymorphism
  - Interface: `signature`
  - Superobject: `std::function<signature>`
  - Subobjects: `int (*)(int)`, `bar`, compatible lambdas
- Interface compliance is checked at **compile time**



```
boost::function_collection<int(int)> c;

std::mt19937 gen{92748}; // some arbitrary random seed
std::discrete_distribution<> rnd{{1,1,1}};
for(int i=0;i<8;++i){ // assign each type with 1/3 probability
    switch(rnd(gen)){
        case 0: c.insert(&foo);break;
        case 1: c.insert(bar{});break;
        case 2: c.insert([](int x){return 3*x;});break;
    }
}

int res=0;
for(const auto& f:c)res+=f(1);
std::cout<<res<<"\n";
```

- Same interface as `boost::base_collection`
- How's this different from `std::vector<std::function<int(int)>>`?

# Duck typing: type erasure on steroids

```
std::ostream& operator<<(std::ostream& os, const sprite& s)
{s.render(os); return os;}
```

```
using concept_=boost::mpl::vector<
    boost::type_erasure::ostreamable<>,
    boost::type_erasure::copy_constructible<>
>;
boost::type_erasure::any<concept_> a=5, b=std::string{"hello"}, c=warrior{16};

std::cout<<a<<"", "<<b<<"", "<<c<<"\n";
```

- Interface: `concept_`
- Superobject: `boost::type_erasure::any<concept_>`
- Subobjects: anything satisfying `concept_`
- Boost.TypeErasure not the only duck typing framework: see for instance Louis Dionne's **Dyno**
- Ultimate solution: native run-time C++ concepts (some day)



```
boost::any_collection<boost::type_erasure::ostreamable<>> c;

std::mt19937 gen{92748}; // some arbitrary random seed
std::discrete_distribution<> rnd{{1,1,1,1,1}};
for(int i=0;i<12;++i){ // assign each type with 1/5 probability
    switch(rnd(gen)){
        case 0: c.insert(warrior{i});break;
        case 1: c.insert(juggernaut{i});break;
        case 2: c.insert(goblin{i});break;
        case 3: c.insert(boost::format{"message %1%"}%i);break;
        case 4: c.insert(i);break;
    }
}

const char* comma="";
for(const auto& x:c){
    std::cout<<comma<<x;
    comma=", ";
}
std::cout<<"\n";
```

- Same interface as before... you get the idea



So many butterflies, so little time



# So many butterflies, so little time

- Boost.PolyCollection packs elements of the same type together for maximum performance
  - Classic OOP      `boost::base_collection`
  - Callable entities    `boost::function_collection`
  - Duck typing      `boost::any_collection`
- Remember: order of elements no longer free
- Squeeze more speed: dedicated algorithms and type restitution
- Dynamic polymorphism goes well beyond OOP
  - Read, play, use
- Drop me a note when you use Boost.PolyCollection



# An intro to Boost.PolyCollection

Thank you

[boost.org/libs/poly\\_collection](http://boost.org/libs/poly_collection)  
[github.com/joaquintides/usingstdcpp2017](https://github.com/joaquintides/usingstdcpp2017)

`using std::cpp 2017`

Joaquín M López Muñoz <[joaquin.lopezmunoz@gmail.com](mailto:joaquin.lopezmunoz@gmail.com)>

Madrid, November 2017